

```
library(dplyr)

rladies_global %>%
  filter(city == 'Auckland')
```



DIVE INTO DPLYR: Manipulating & exploring your data



Hello!

Welcome to R-Ladies



1.

Introduction

R language, RStudio,
tidyverse



Three things you'll need to install

1. **Install R** -- this is the open-source programming language we'll use (download via CRAN -- Comprehensive R Archive Network)
2. **Install RStudio** -- this is the most popular IDE for R and will make your life a lot easier (download from rstudio.com/download)
3. **Install dplyr** -- this is the package we'll use within R to work with data. Install with one line of code in R:
`install.packages("dplyr")`



What is the tidyverse?

- Collection of R packages based on tidy data principles
- Designed to work together
- An easier way to code!
- AKA “Hadleyverse” (most packages written by Hadley Wickham)
- Dplyr is one of the main packages

What is the tidyverse?



What is tidy data?


- Each variable is a column
- Each observation is a row
- Each type of observational unit is a table

id	artist	track	time
1	2 Pac	Baby Don't Cry	4:22
2	2Ge+her	The Hardest Part Of ...	3:15
3	3 Doors Down	Kryptonite	3:53
4	3 Doors Down	Loser	4:24
5	504 Boyz	Wobble Wobble	3:35
6	98~0	Give Me Just One Nig...	3:24
7	A*Teens	Dancing Queen	3:44
8	Aaliyah	I Don't Wanna	4:15
9	Aaliyah	Try Again	4:03
10	Adams, Yolanda	Open My Heart	5:30
11	Adkins, Trace	More	3:05
12	Aguilera, Christina	Come On Over Baby	3:38
13	Aguilera, Christina	I Turn To You	4:00
14	Aguilera, Christina	What A Girl Wants	3:18
15	Alice DeeJay	Better Off Alone	6:50



2. dplyr

Let's get started!





What is dplyr?

Grammar of data manipulation:

- + **select()** picks variables (columns) based on their names
- + **arrange()** changes the ordering of rows
- + **filter()** allows row selection based on given criteria
- + **mutate()** creates new variables (columns) from existing ones
- + **summarise()** reduces multiple values down to a single summary
- + **group_by()** performs any of the above on a group-by-group basis

dplyr

syntax

- + All calls to dplyr verbs follow the same format:
 1. The first argument is a dataframe
 2. The subsequent arguments describe what to do to that dataframe, using unquoted variable names.
- + Each call returns a new dataframe (rather than overwriting the 'old' one)
- + Example:
`filter(.data = iris, Species == "setosa")`

Quick aside: iris dataset

- + Included in R (**iris** to view)
- + 150 observations of 5 variables:
Iris type, sepal length + width, and petal length + width



select()

- + Picks variables (columns) based on their names
- + First argument is dataframe; subsequent arguments represent columns to select

select(.data = iris , Species, Petal.Length, Petal.Width)

```
> iris
  Sepal.Length Sepal.Width Petal.Length Petal.Width Species
1          5.1         3.5          1.4          0.2   setosa
2          4.9         3.0          1.4          0.2   setosa
3          4.7         3.2          1.3          0.2   setosa
4          4.6         3.1          1.5          0.2   setosa
5          5.0         3.6          1.4          0.2   setosa
```

```
Species Petal.Length Petal.Width
1   setosa          1.4          0.2
2   setosa          1.4          0.2
3   setosa          1.3          0.2
4   setosa          1.5          0.2
5   setosa          1.4          0.2
```

select() + helper functions

Helper functions you can use within `select()`:

- + `starts_with("a")` matches names that begin with "a"
- + `ends_with("z")` matches names that begin with "z"
- + `contains("lady")` matches names that contain "lady"
- + `matches(<regex>)` allows you to do regex matching on names

```
> select(iris, starts_with("P"))  
  Petal.Length Petal.Width  
1          1.4          0.2  
2          1.4          0.2  
3          1.3          0.2  
4          1.5          0.2  
5          1.4          0.2  
6          1.7          0.4
```



You do it!

- + Select the first 3 columns of iris
- + Select Petal.Width and Petal.Length
- + Select columns with Sepal in them
- + Select all columns BUT Species

arrange()

- + Changes the ordering of rows
- + First argument is the dataframe, subsequent arguments are columns and/or expressions used to re-arrange the dataframe
- + Note: default is ascending order, and NA's are always at the end

`arrange(.data = iris, Sepal.Length, Sepal.Width)`

```
> iris
  Sepal.Length Sepal.Width Petal.Length Petal.Width Species
1         5.1         3.5          1.4          0.2   setosa
2         4.9         3.0          1.4          0.2   setosa
3         4.7         3.2          1.3          0.2   setosa
4         4.6         3.1          1.5          0.2   setosa
5         5.0         3.6          1.4          0.2   setosa
```

```
  Sepal.Length Sepal.Width Petal.Length Petal.Width Species
1         4.3         3.0          1.1          0.1   setosa
2         4.4         2.9          1.4          0.2   setosa
3         4.4         3.0          1.3          0.2   setosa
4         4.4         3.2          1.3          0.2   setosa
5         4.5         2.3          1.3          0.3   setosa
```



You do it!

- + Sort iris by column Petal.Width
- + Sort by Petal.Width descending
- + Sort by Petal.Width and then Petal.Length

filter()

- + Allows pointed row selection based on given criteria
- + First argument is the dataframe, subsequent arguments are logical expressions used to filter the dataframe

```
filter( .data = iris, Species == "setosa")
```

```
> iris
  Sepal.Length Sepal.Width Petal.Length Petal.Width Species
1          5.1         3.5          1.4          0.2  setosa
2          4.9         3.0          1.4          0.2  setosa
3          4.7         3.2          1.3          0.2  setosa
4          4.6         3.1          1.5          0.2  setosa
5          5.0         3.6          1.4          0.2  setosa
```

nrow = 150

```
  Sepal.Length Sepal.Width Petal.Length Petal.Width Species
1          5.1         3.5          1.4          0.2  setosa
2          4.9         3.0          1.4          0.2  setosa
3          4.7         3.2          1.3          0.2  setosa
4          4.6         3.1          1.5          0.2  setosa
5          5.0         3.6          1.4          0.2  setosa
```

nrow = 50

filter() + booleans

Multiple arguments to `filter()` are combined with “and”: every expression must be true in order for a row to be included in the output. For other types of combinations, you’ll need to use Boolean operators yourself: `&` is “and”, `|` is “or”, and `!` is “not”. Figure 5.1 shows the complete set of Boolean operations.




Figure 5.1: Complete set of boolean operations. x is the left-hand circle, y is the right-hand circle, and the shaded region show which parts each operator selects.



You do it!

- + Find the rows where `Petal.Width > 1`
- + Find the rows where `Petal.Width > 1` and `Species` is `versicolor`
- + Find the rows where `Species` is not `setosa`

3. magrittr

What?!?



What is magrittr?

Simplifying R code with pipes (`%>%`)

- + Easy way to pass data through functions without nesting
- + First argument of each function is “piped” in to reduce redundancy
- + $f(x)$ is the same as `x %>% f()`
- + $f(x, y)$ is the same as `x %>% f(y)`



dplyr + magrittr example

Before

```
> select(  
+   filter(iris, Petal.Width > 1),  
+   Species, Petal.Length)
```

After

```
> iris %>%  
+   filter(Petal.Width > 1) %>%  
+   select(Species, Petal.Length)
```



You do it!

- + Find the rows where `Petal.Width > 1` and select the first 3 variables
- + Select the columns that contain `Petal` and sort by `Petal.Width` and then only the first 6 rows
- + Find the rows for Species `setosa` and order rows by descending `Sepal.Width`
- + Select the rows where `Petal.Width > 1` and view in the panel
- + Find the unique values of `Petal.Width` for the `setosa` Species

Quick aside:

Missing values

- + NA represents a missing (unknown) value
- + Comparisons involve unknown values typically result in unknown values
- + To see whether a value is missing, use **is.na()**
- + **filter()** only includes rows where the condition is true (not false or NA)

```
# Let x be Mary's age. We don't know how old she is.  
x <- NA  
  
# Let y be John's age. We don't know how old he is.  
y <- NA  
  
# Are John and Mary the same age?  
x == y  
#> [1] NA  
# We don't know!
```


mutate()

- + Creates new variables (columns) from existing ones
- + Note: columns created with `mutate()` are always added to end of dataset

```
iris %>% mutate( petal_area = Petal.Length * Petal.Width)
```

```
> iris
  Sepal.Length Sepal.Width Petal.Length Petal.Width Species
1          5.1         3.5          1.4          0.2  setosa
2          4.9         3.0          1.4          0.2  setosa
3          4.7         3.2          1.3          0.2  setosa
4          4.6         3.1          1.5          0.2  setosa
5          5.0         3.6          1.4          0.2  setosa
```

```
  Sepal.Length Sepal.Width Petal.Length Petal.Width Species petal_area
1          5.1         3.5          1.4          0.2  setosa      0.28
2          4.9         3.0          1.4          0.2  setosa      0.28
3          4.7         3.2          1.3          0.2  setosa      0.26
4          4.6         3.1          1.5          0.2  setosa      0.30
5          5.0         3.6          1.4          0.2  setosa      0.28
```

mutate()

Useful functions

- + Arithmetic operators (+, -, *, /, ^)
- + Log functions (like `log10()`)
- + Offsets like `lead()` and `lag()`
- + Logical comparisons (<, <=, >, >=, !=)
- + `ifelse` statements (if this, then this, else this)
- + Or when more than 1 logical split then use `case_when`
- + Cumulative and rolling aggregates
- + Ranking (like `ntile()`)



You do it!

- + Create a new column with a flag showing $\text{Petal.Width} > \text{Petal.Length}$
- + Create a new column with $\text{Petal.Width} + \text{Sepal.Length}$
- + Create a new column with the mean of Sepal.Length and then another new column with a flag stating if each row's Sepal.Length is greater than the mean
- + Create a new column with Sepal.Length buckets (use `cut`)
- + Create a new column stating Low when $\text{Petal.Width} < 0.2$, Medium when Petal.Width is between 0.2 and 0.6 and High when $\text{Petal.width} > 0.6$ (use `case_when`)

group_by() and summarise()

- + group_by applies dplyr verbs by group
- + summarise reduces multiple values down to a single summary

```
iris %>%  
  group_by(Species) %>%  
  summarise(avg_petal_width = mean(Petal.Width))
```

```
> iris
```

	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
1	5.1	3.5	1.4	0.2	setosa
2	4.9	3.0	1.4	0.2	setosa
3	4.7	3.2	1.3	0.2	setosa
4	4.6	3.1	1.5	0.2	setosa
5	5.0	3.6	1.4	0.2	setosa

```
# A tibble: 3 × 2  
  Species avg_petal_width  
  <fctr>      <dbl>  
1   setosa      0.246  
2 versicolor  1.326  
3  virginica   2.026
```



summarise()

Useful functions

- + Counts (**n()**, **n_distinct()**)
- + Measures of location (**mean()**, **median()**)
- + Measures of spread (**sd()**, **IQR()**)
- + Measures of position (**first()**, **last()**)

You do it!

- + Find the mean of the Petal.Length and the sd of Petal.Width
- + Create a new column with the mean of Sepal.Length per species
- + Create a table with the mean Sepal.Length per species
- + Find how many observations there are for each species and Petal.Length combination

Tips & Tricks

- + If you don't have the result of a dplyr chain to a dataframe, it will print
- + Include arguments for inner functions like `na.rm = TRUE`
- + `rename()` is a cool function to clean up messy column names
- + `slice()` will give you rows by row number
- + After grouping with `group_by()`, you can `ungroup()` to remove groupings
- + There is a [cheat sheet](#) for data wrangling!
- + If you want to print and save, wrap assignment in parenthesis
 - + Example: `(iris_names <- iris %>% filter(Species == "setosa"))`



4.

Scoped verbs

More dplyr goodness . . .



Scoped Verbs

- + Terminology: we have been using “single table verbs”
- + Now we can affect multiple variables simultaneously with the scoped verbs
- + Three extensions
 - + `_if` pick variables based on a predicate function like `is.numeric()`
 - + `_at` pick variables using the same syntax as `select()`.
 - + `_all` operates on all variables

summarise_all()

Scoped verbs

- + apply the function to all variables

`summarise_all(.tbl, function)`

- + Multiple functions `summarise_all(.tbl, funs(f1, f2))`
- + Combine functions: `summarise_all(.tbl, ~f1(f2(.)))`

```
> iris %>%  
+   summarise_all(length)  
  Sepal.Length Sepal.Width Petal.Length Petal.Width Species  
1           150          150          150          150      150
```



You do it!

- + Find the mean of the first 4 variables
- + Find the mean and min of the first 4 variables
- + Find how many unique observations there are for each of the variables



summarise_at()

Scoped verbs

- + apply the function to variables using their names like select

```
summarise_at( .tbl, vars(...), function)
```

```
> iris %>%  
+   summarise_at(vars(-Species), mean)  
  Sepal.Length Sepal.Width Petal.Length Petal.Width  
1      5.843333      3.057333      3.758      1.199333
```



You do it!

- + Find the mean of the first 4 variables
- + Find the mean and min of the variables with Petal in their names

summarise_if()

Scoped verbs

- + apply the function to variables using logical conditions based on some property of the column

`summarise_if(.tbl, condition, function)`

```
> iris %>%  
+   summarise_if(is.numeric, mean)  
  Sepal.Length Sepal.Width Petal.Length Petal.Width  
1      5.843333      3.057333         3.758      1.199333
```



You do it!

- + Find the sd of the numerical variables
- + Find the number of unique observations of all the categorical variables



mutate & filter

My frequent examples

- + `mutate_if(.tbl, is.character, as.factor)`
- + `filter_all()` used with `all_vars()` or `any_vars()`
- + `filter_all(data, any_vars(is.na(.)))`



You do it!

- + Turn numerical variables into characters
- + Round the Petal variable to 0dp
- + Find the rows of the airquality dataset where there is at least 1 NA value



5. Wrap-up

Announcements, upcoming events, etc.



R-Ladies Auckland

Upcoming Events

Lessons and highlights from the rOpenSci project [Feb 20]

useR! [Brizzie Jul 10-13]

Looking for presenters!



BIG THANKS TO

R-Ladies Austin

https://github.com/rladiesaustin/R4DS_workshop_series

Hadley Wickham vignettes and Stanford material

<https://dcl-2017-04.github.io/curriculum/manip-scoped.html>